

# Yashwanth Reddy Boddireddy

Data & AI Professional | ML, GenAI, and Scalable Analytics, Jersey City, NJ  
+1 (201) 884-3230 | [yashwanthreddy335@gmail.com](mailto:yashwanthreddy335@gmail.com) | [linkedin.com/in/ yashwanth-reddy-boddireddy](https://www.linkedin.com/in/yashwanth-reddy-boddireddy) |  
[GitHub.com/yashwanthreddy7178](https://github.com/yashwanthreddy7178)

## SUMMARY

Senior Data & AI professional with 5 years of progressive experience designing and deploying large-scale data pipelines, machine learning systems, and generative AI solutions. Proven record of delivering 3 TB+ daily ETL pipelines, 50K+ inferences per day, and advanced analytics that boosted user retention by 15%, accelerated model rollout by 80%, and drove \$5M+ in product investments.

## SKILLS

**Programming & Frameworks:** Python, SQL, Tensorflow, Pytorch, Scikitlearn, Spark

**Machine Learning & GenAI:** AI Agents, Langchain, FastAPI, Anomaly Detection, Preductive Modelling

**MLOps & DevOps:** Docker, Kubernetes, Jenkins, MLflow, Databricks, Flask, Snowflake

**Cloud Platforms & Tools:** AWS(Lambda, S3, Glue, RDS, EC2 GPU), Azure((Data Factory) ,GCP (BigQuery)

**Data Engineering & Orchestration:** Airflow, Dask, Hadoop, dbt, ETL Pipelines, Data Modeling, MongoDB, kafka

**Collaboration & Leadership:** Agile, Cross-functional Teams, Stakeholder Demos, Mentoring, Documentation

## EXPERIENCE

### Senior Data Scientist

Vmware, Remote, USA

July 2024 – Present

- Engineered a scalable MLOps platform (Docker, Kubernetes, Jenkins, Airflow) delivering 50K+ real-time inferences/day with 80% faster rollouts, zero-downtime CI/CD, and less than 5% prediction error.
- Boosted ML performance by 22% F1 via large-scale feature engineering (SageMaker, EMR, Dask, MLflow) and achieved 94% accuracy in time-series models and 28% higher precision in anomaly detection.
- Automated 120+ ML pipelines across 500+ systems (Airflow, Databricks) and containerized 30+ microservices, processing 3 TB/day of telemetry to accelerate training and iteration.
- Delivered RAG-based LLM assistant and API solutions (OpenAI, LangChain), parsed 1 TB+ logs with TensorFlow anomaly detection, and created dashboards that drove 15% retention lift and informed \$5 M+ investments.

### Senior Data Scientist

Accenture, Hyderabad, India

October 2020 – August 2023

- Built HBO Max's real-time recommendation engine using collaborative filtering and NLP, boosting session time by 19%, MAUs by 12%, and trailer conversions by 30%.
- Migrated Spark pipelines to Docker-on-ECS with Spinnaker CI/CD and unified 5 TB/day streaming into an S3 data lake via Airflow + AWS Glue, cutting model delivery time by 30%.
- Implemented CloudWatch/Prometheus monitoring, custom A/B and multi-armed bandit frameworks, reducing churn incidents by 40% and improving YoY churn by 8.5% and CTR by 2%.
- Modularized ML assets in Databricks with MLflow lineage, mentored three junior data scientists, and led tech talks, doubling team velocity and elevating code quality.

## PROJECTS

### LLM-Powered Root Cause Explorer for Anomaly Investigation

- Designed and deployed an LLM-powered RCA system (LangChain + DistilBERT + RAG) on AWS ECS Fargate with Kafka + Prometheus integration, achieving 87% precision across 10+ anomaly types and reducing triage time by 65%.

### Scalable Credit Risk Scoring System

- Built a real-time credit scoring system with PySpark + XGBoost (AWS Glue + Kubernetes), achieving 89% F1-score, sub-100ms latency, and regulatory-aligned explainability via SHAP dashboards, boosting stakeholder trust and compliance.

## EDUCATION

Master of Science in Data Science, Statistics, New Jersey Institute of Technology, USA

September 2023 – May 2025